US006137043A

United States Patent [19]

Rossum

[11] **Patent Number: 6,137,043**

[45] **Date of Patent: Oct. 24, 2000**

[54] **DIGITAL SAMPLING INSTRUMENT EMPLOYING CACHE MEMORY**

[75] Inventor: **David P. Rossum**, Aptos, Calif.

[73] Assignee: **Creative Technology Ltd.**, Singapore, Singapore

[21] Appl. No.: **09/187,139**

[22] Filed: **Nov. 6, 1998**

Related U.S. Application Data

[60] Continuation of application No. 08/903,329, Jul. 29, 1997, Pat. No. 5,925,841, which is a division of application No. 08/636,827, Apr. 23, 1996, Pat. No. 5,698,803, which is a continuation of application No. 08/202,922, Feb. 28, 1994, abandoned, which is a division of application No. 07/882,178, May 11, 1992, Pat. No. 5,342,990, which is a continuation-in-part of application No. 07/462,392, Jan. 5, 1990, Pat. No. 5,111,727.

[51] **Int. Cl.**[7] **G10H 7/10**; G10H 7/12

[52] **U.S. Cl.** **84/603**; 84/604; 84/607; 708/290

[58] **Field of Search** ..... 84/601–608, 623, 84/659–661, 693, DIG. 9, DIG. 30; 364/134, 723, 724.1, 724.12, 728.01, 728.02; 708/290, 313, 315, 420, 421

[56] **References Cited**

U.S. PATENT DOCUMENTS

Re. 32,862	2/1989	Wachi	84/1.21
3,988,607	10/1976	Eggermont et al.	235/156
3,997,773	12/1976	Van Essen et al.	235/156
4,020,332	4/1977	Crochiere et al.	235/152
4,231,277	11/1980	Wachi	84/1.19
4,246,823	1/1981	Wachi et al.	84/1.22
4,377,960	3/1983	Okumura	84/1.01
4,432,265	2/1984	Oya et al.	84/1.01
4,460,890	7/1984	Busby	340/347
4,643,067	2/1987	Deutsch	84/1.01
4,702,142	10/1987	Deutsch	84/1.01
4,715,257	12/1987	Hoshiai et al.	84/1.01
4,829,463	5/1989	Kakishita et al.	364/724.19
4,901,615	2/1990	Matsushima et al.	84/605
5,007,323	4/1991	Usami	84/607
5,023,825	6/1991	Luthra et al.	364/724.1
5,067,141	11/1991	Critchlow et al.	375/97
5,073,942	12/1991	Yoshida et al.	381/61
5,111,727	5/1992	Rossum	84/603
5,245,667	9/1993	Lew	381/94
5,300,724	4/1994	Medovich	84/604
5,342,990	8/1994	Rossum	84/603
5,698,803	12/1997	Rossum	84/607
5,925,841	7/1999	Rossum	84/603

OTHER PUBLICATIONS

The Technology of Computer Music, Mathews et al., 1969.

Multirate Digital Signal Processing, Crochiere & Rabiner, 1983.

Introduction to Numerical Analysis, Hildebrand, 1956.

Digidesign's Sound Accelerator: Lessons Lived and Learned, Bill Lowe et al., Computer Musical Journal, vol. 13, No. 1, 1989.

James A. Moorer et al., *"The Digital Audio Processing Station: A New Concept in Audio Postproduction,"* 78th Convention of Audio Engineering Society, Anaheim, CA, May 1985.

"Kurzweil 1000 Series," Service Manual (1000 Racks, K1000, EGP and Mark III Models), Kurzweil Music Systems, Inc. Waltham, MA, Jan. 1989, Doc. No.92000101.

Musician's Guide, Ralph Jones et al., K1000SE, Kurzweil Music Systems, Inc., Waltham, MA, 1988.

AKAI professional "Service Bulletin/HD80 Non–Operation of Fan", Dec. 12, 1990.

AKAI Service Manual, Model S10000KB, Mar. 15, 1990.

Sequential/USA Model 2000—Prophet 2000—Digital Sampling Keyboard Technical Manual, Rick Davies, Dec. 1985.

(List continued on next page.)

Primary Examiner—Stanley J. Witkowski

Attorney, Agent, or Firm—Townsend and Townsend and Crew LLP

[57] **ABSTRACT**

A digital sampling instrument for multi-channel interpolatative playback of digital audio data stored in a waveform memory provides improved interpolation of musical sounds by use of a cache memory.

12 Claims, 12 Drawing Sheets

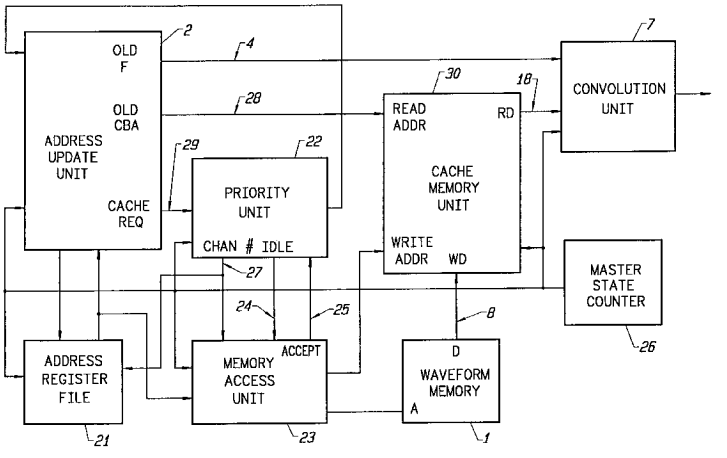

OTHER PUBLICATIONS

E–mu Systems Emulator—Operating Instructions, Marco Alpert et al., Version 3.6, 1981.
E–mu Systems—Emulator Service Manual—Revision of Jan. 12, 1982.
Sequential/Europe—Prophet 2000 Digital Sampling Keyboard and Prophet 2002 Rack–Mount Sampler Operation Manual, Stanley Jungleib, Apr. 1986.
EMAX Digital Sampler Technical Manual, Riley B. Smith, 1987, E–mu Sysetms Inc.
E–mu Systems Inc., Emax HD Digital Sampling Keyboard, Owner's Manual, Craig Anderson, 1988.
Dancetech—Casio FZ Range . . . FZI . . . FZ–10M . . . FZ–20M, Aug. 25, 1999.
Computer Storage Systems & Technology, Richard E. Matick, 1977.
ICASSP 84—Proceedings—IEEE International Conference on Acoustics, Speech, and Signal Processing, vol. 2 of 3.
Some Aspects of Sample Rate Conversion, Dave Rossum, E–mu Systems, ICMC '85 Proceedings.
An Analysis of Pitch Shifting Algorithms, Dave Rossum, E–mu Systems, AES 87th Convention, Oct. 1989.
Compiling the E–Chip Music Signal Processor, Dave Rossum, E–mu Systems, Automated Design & Engineering for Electronics West, 1986.
Compiling a Music Signal Processor, Dave Rossum, E–mu Systems, Inc., 1986 IEEE.
Practical Considerations in the Design of Music Systems Using VLSI, J. William Mauchly et al., AES 5th International Conference 1987.
Fourier transform and convolution subroutines for the IBM 3090 Vector Facility, Ramash c. Agarwal et al., IBM J. Res. Develop. vol. 30, No. 2, Mar. 1986.
Accelerando: A Real–Time, General Purpose Computer Music System, Keith Lent et al., Computer Music Journal, vol. 13, No. 4, Winter 1989.
An Efficient Method for Pitch Shifting Digitally Sampled Sounds, Keith Lent, Computer Music Journal, vol. 13, No. 4, Winter 1989.
An Overview of the Sound and Music Kits for the NeXT Computer, David Jaffe et al., Computer Music Journal, vol. 13, No. 2, Summer 1989.

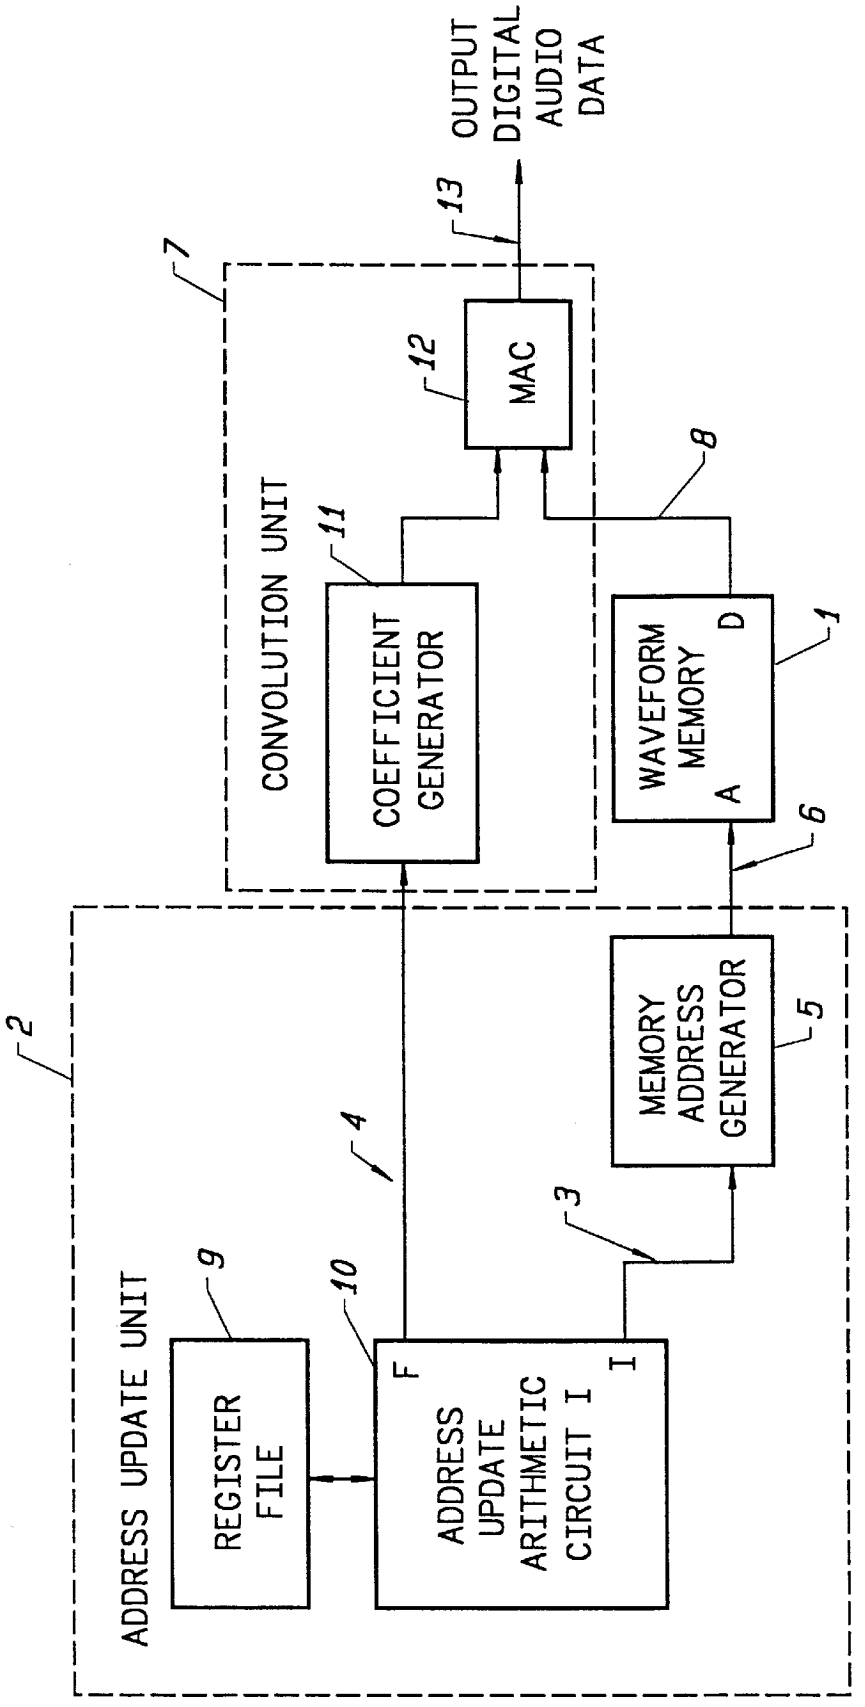
OUTPUT DIGITAL AUDIO DATA
13
7
CONVOLUTION UNIT
12
MAC
8
11
COEFFICIENT GENERATOR
WAVEFORM MEMORY
D
A
1
6
MEMORY ADDRESS GENERATOR
5
2
4
3
9
10
ADDRESS UPDATE UNIT
REGISTER FILE
F
ADDRESS UPDATE ARITHMETIC CIRCUIT I
I

FIG. 1

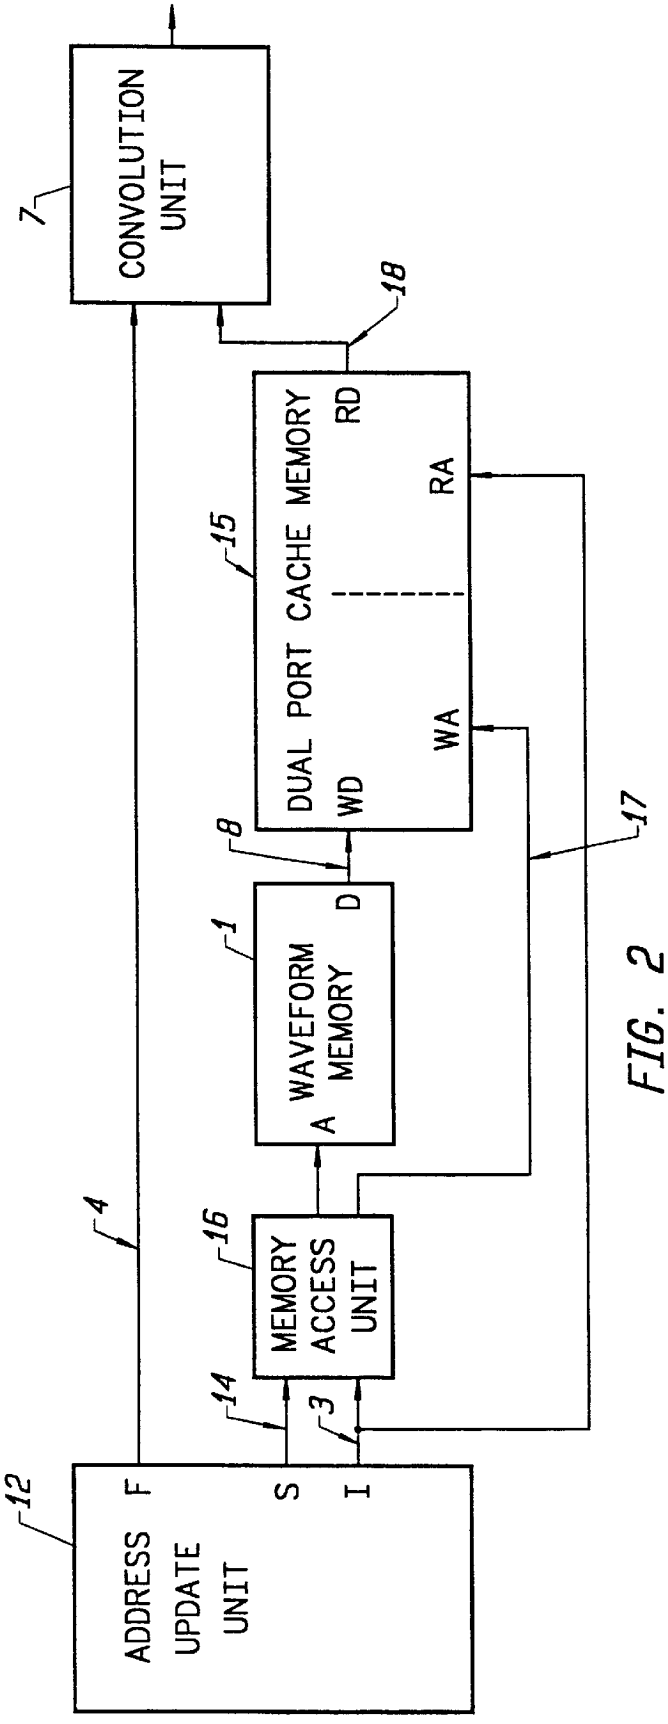
CONVOLUTION UNIT
7
18
DUAL PORT CACHE MEMORY
RD
RA
WA
WD
15
8
17
WAVEFORM MEMORY
A
D
1
4
MEMORY ACCESS UNIT
16
14
3
ADDRESS UPDATE UNIT
12
F
S
I

FIG. 2

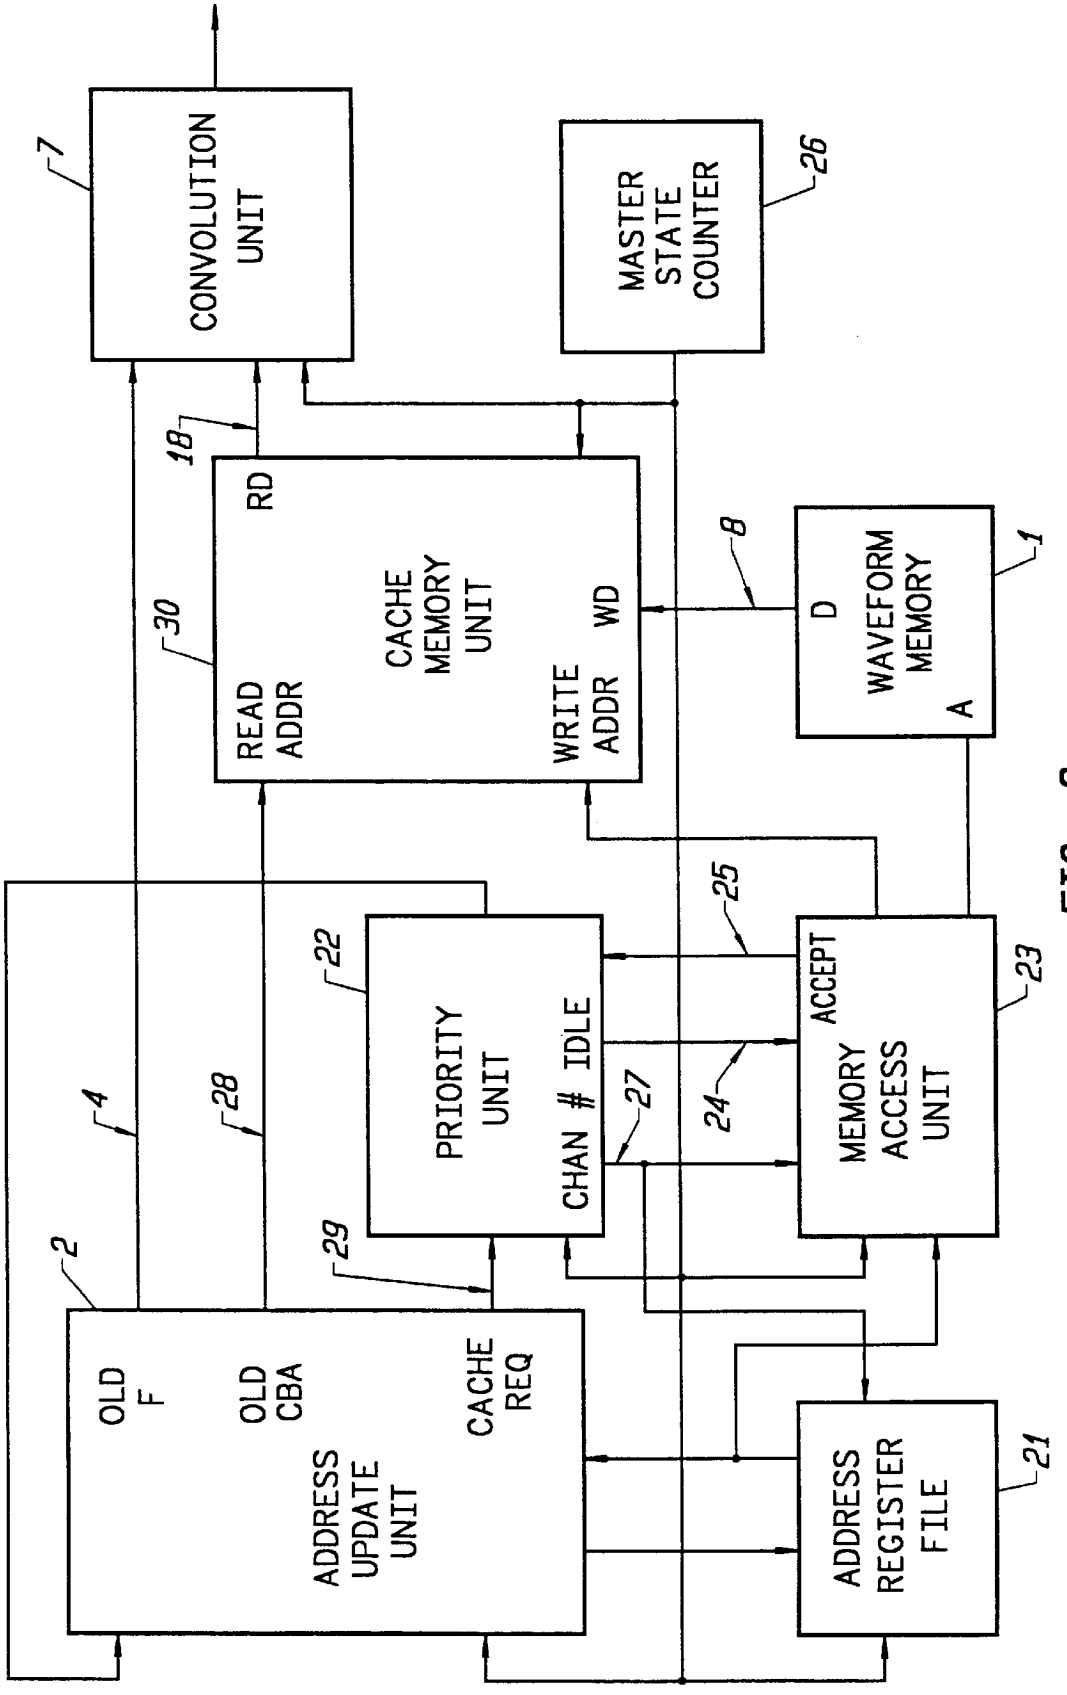
CONVOLUTION UNIT
7
MASTER STATE COUNTER
26
18
RD
CACHE MEMORY UNIT
WD
30
READ ADDR
WRITE ADDR
8
D
WAVEFORM MEMORY
A
1
25
22
PRIORITY UNIT
CHAN # IDLE
27
24
ACCEPT
MEMORY ACCESS UNIT
23
4
28
29
2
OLD F
OLD CBA
CACHE REQ
ADDRESS UPDATE UNIT
ADDRESS REGISTER FILE
21

FIG. 3

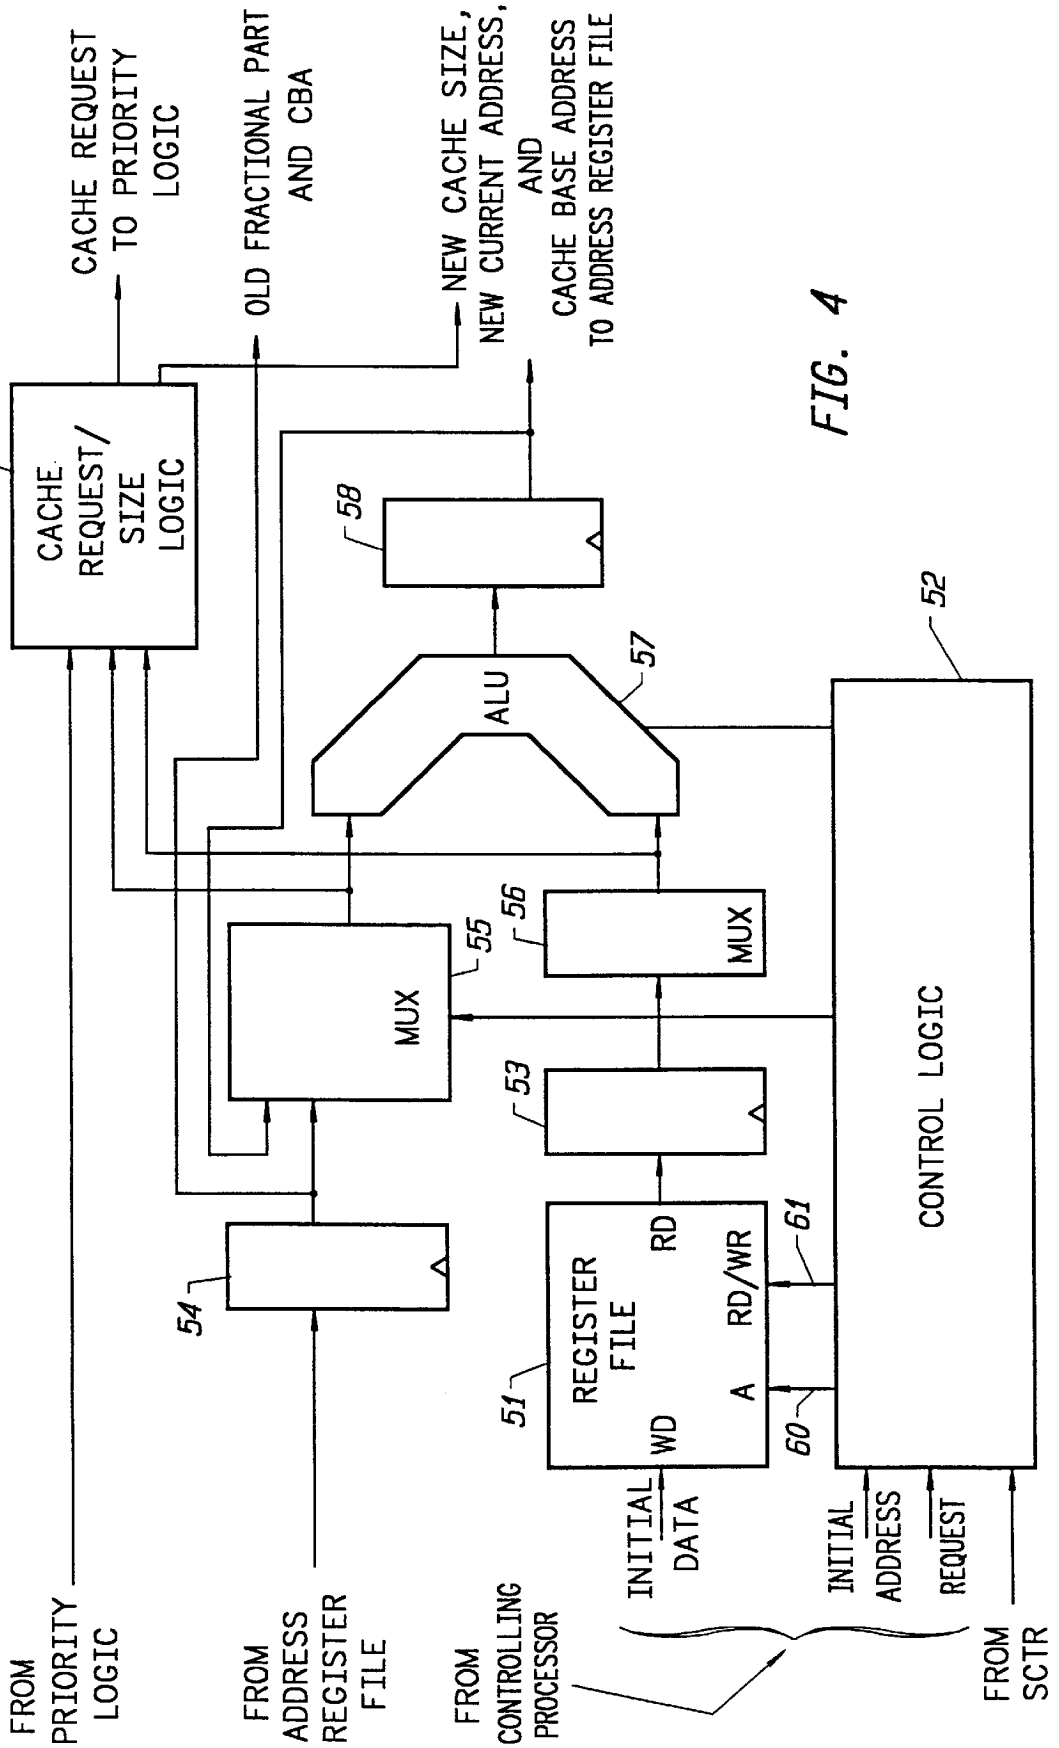
CACHE REQUEST TO PRIORITY LOGIC
OLD FRACTIONAL PART AND CBA
NEW CACHE SIZE, NEW CURRENT ADDRESS, AND CACHE BASE ADDRESS TO ADDRESS REGISTER FILE
59
CACHE REQUEST/ SIZE LOGIC
58
ALU
57
52
MUX
55
56
MUX
53
54
REGISTER FILE
WD
RD
RD/WR
A
51
60
61
CONTROL LOGIC
FROM PRIORITY LOGIC
FROM ADDRESS REGISTER FILE
FROM CONTROLLING PROCESSOR
INITIAL DATA
INITIAL ADDRESS
REQUEST
FROM SCTR

FIG. 4

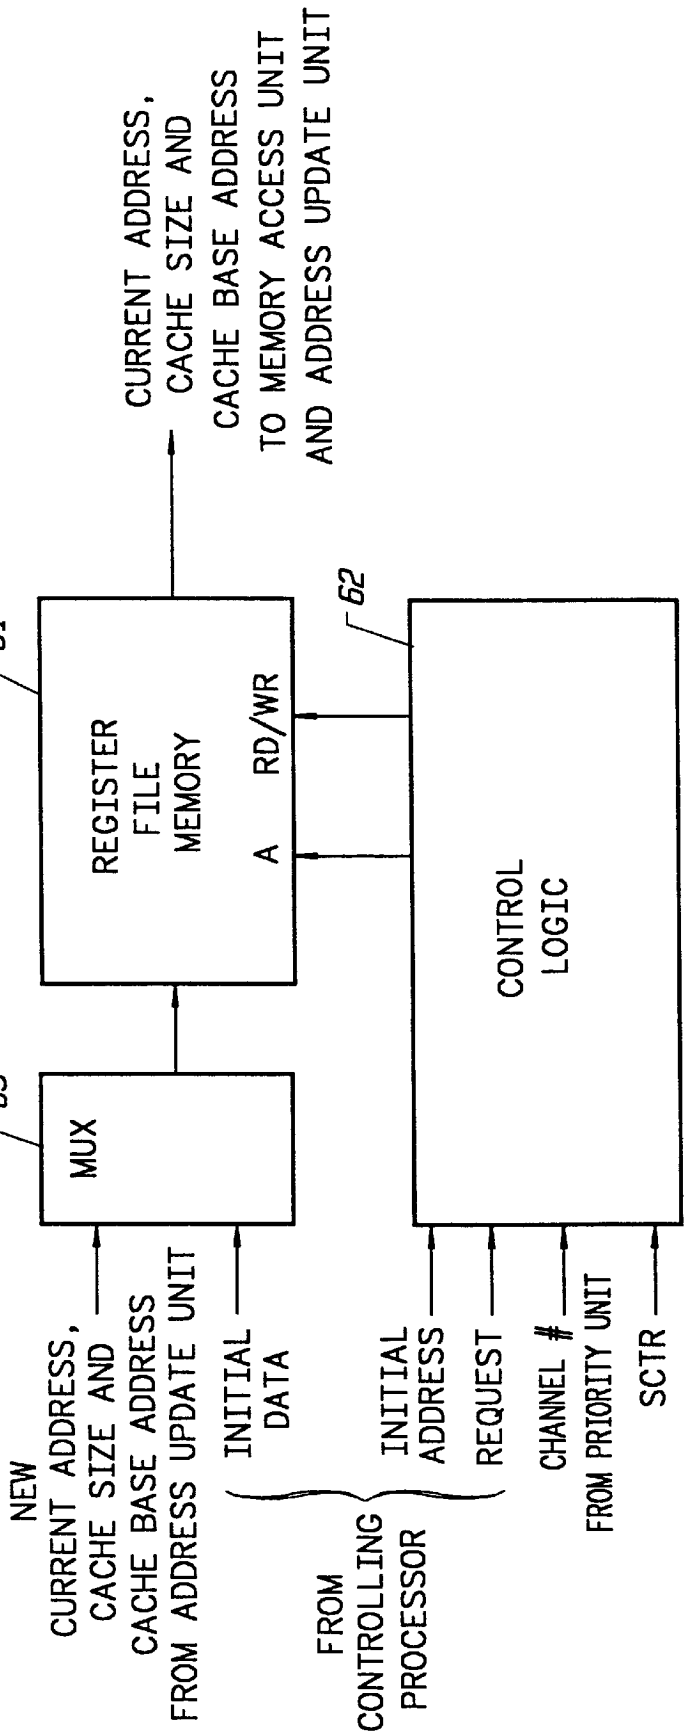
61
REGISTER FILE MEMORY
A
RD/WR
62
CONTROL LOGIC
63
MUX
CURRENT ADDRESS, CACHE SIZE AND CACHE BASE ADDRESS TO MEMORY ACCESS UNIT AND ADDRESS UPDATE UNIT
NEW CURRENT ADDRESS, CACHE SIZE AND CACHE BASE ADDRESS FROM ADDRESS UPDATE UNIT
INITIAL DATA
FROM CONTROLLING PROCESSOR
INITIAL ADDRESS
REQUEST
CHANNEL # FROM PRIORITY UNIT
SCTR

FIG. 5

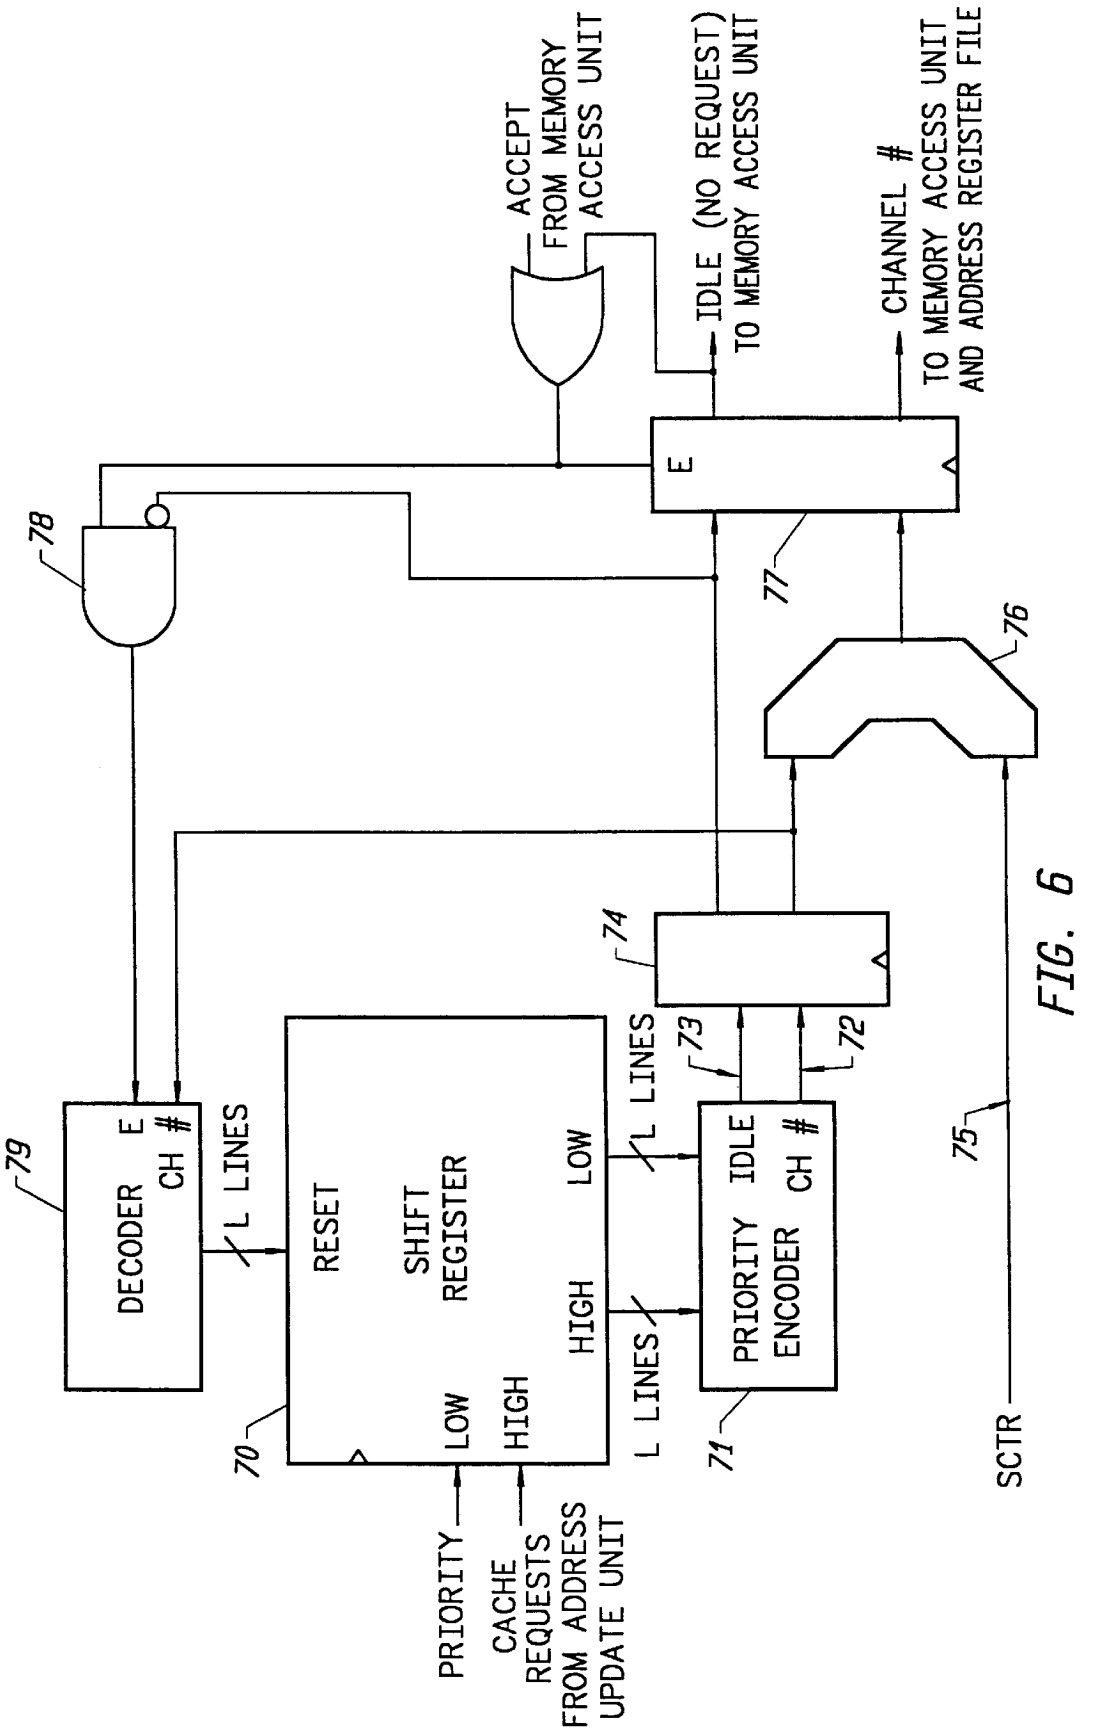
ACCEPT FROM MEMORY ACCESS UNIT
IDLE (NO REQUEST) TO MEMORY ACCESS UNIT
CHANNEL # TO MEMORY ACCESS UNIT AND ADDRESS REGISTER FILE
E
78
77
76
74
73
72
75
79
DECODER E CH #
L LINES
RESET
SHIFT REGISTER
LOW
HIGH
L LINES
L LINES
PRIORITY IDLE ENCODER CH #
70
71
PRIORITY
CACHE REQUESTS FROM ADDRESS UPDATE UNIT
SCTR

FIG. 6

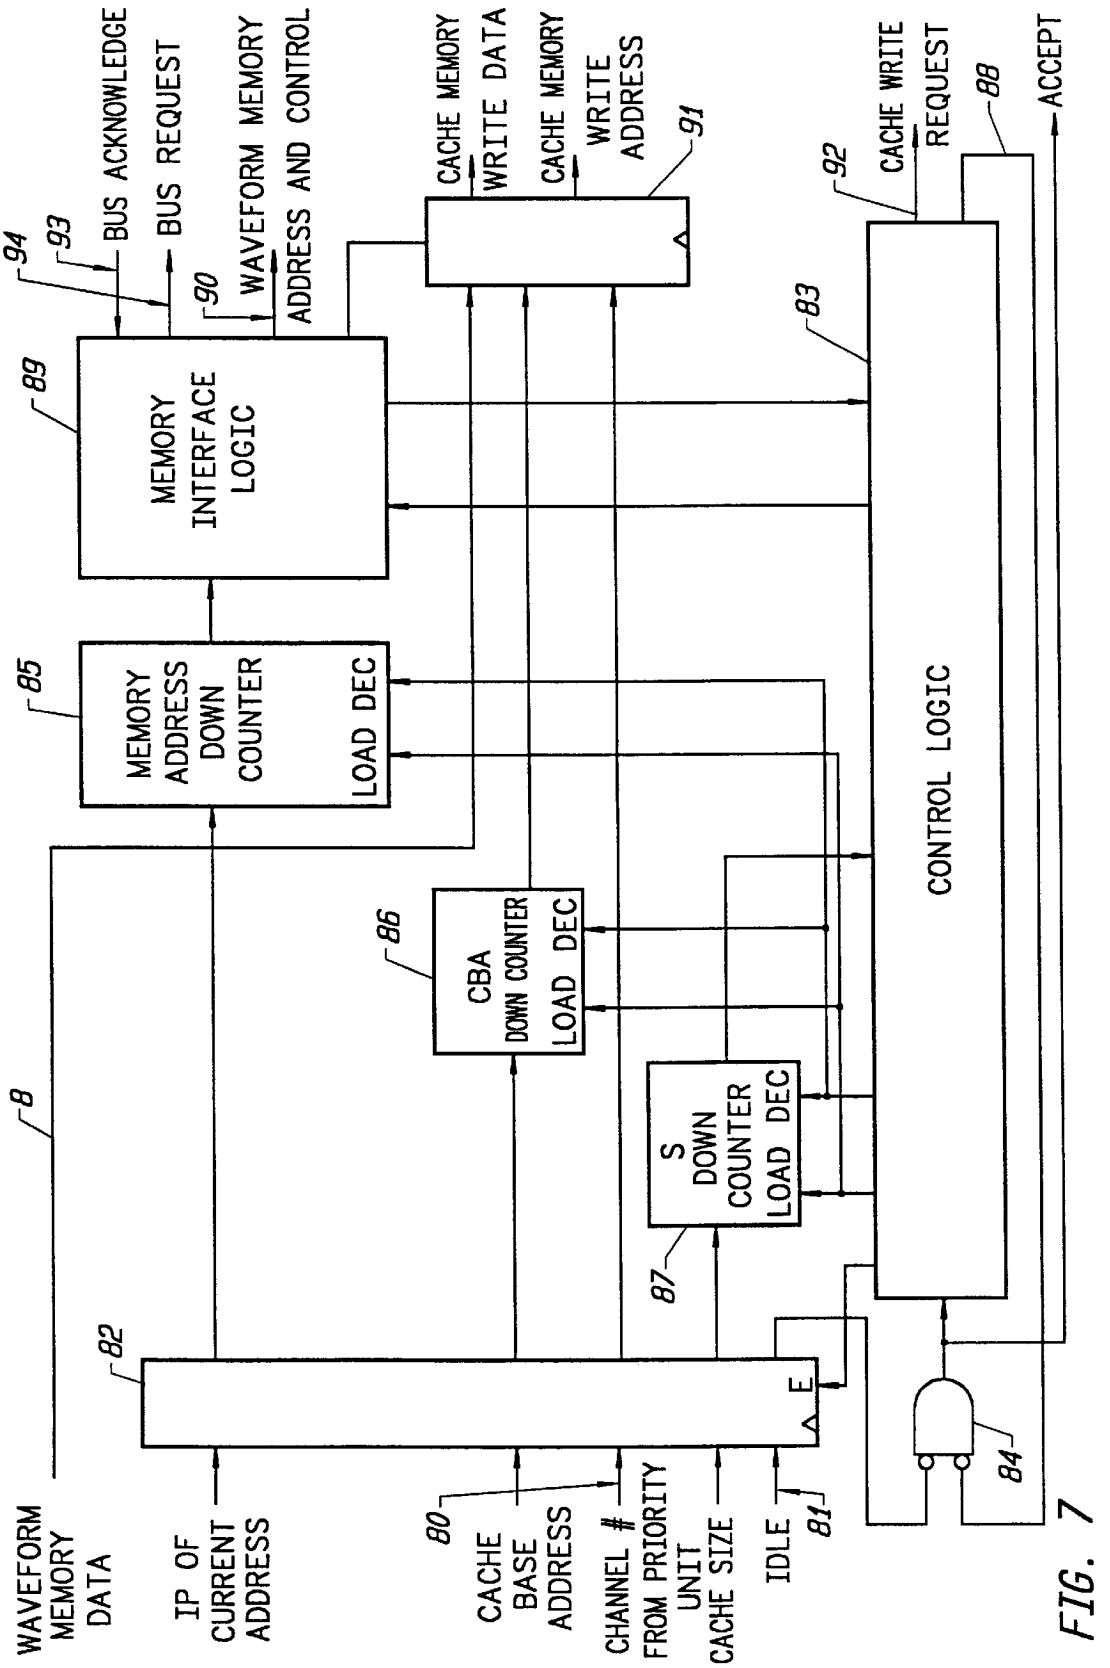
WAVEFORM MEMORY DATA
IP OF CURRENT ADDRESS
CACHE BASE ADDRESS
CHANNEL # FROM PRIORITY UNIT
CACHE SIZE
IDLE
BUS ACKNOWLEDGE
BUS REQUEST
WAVEFORM MEMORY ADDRESS AND CONTROL
CACHE MEMORY WRITE DATA
CACHE MEMORY WRITE ADDRESS
CACHE WRITE REQUEST
ACCEPT
MEMORY INTERFACE LOGIC
MEMORY ADDRESS DOWN COUNTER
LOAD DEC
CBA DOWN COUNTER
LOAD DEC
S DOWN COUNTER
LOAD DEC
CONTROL LOGIC
E
8
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94

FIG. 7

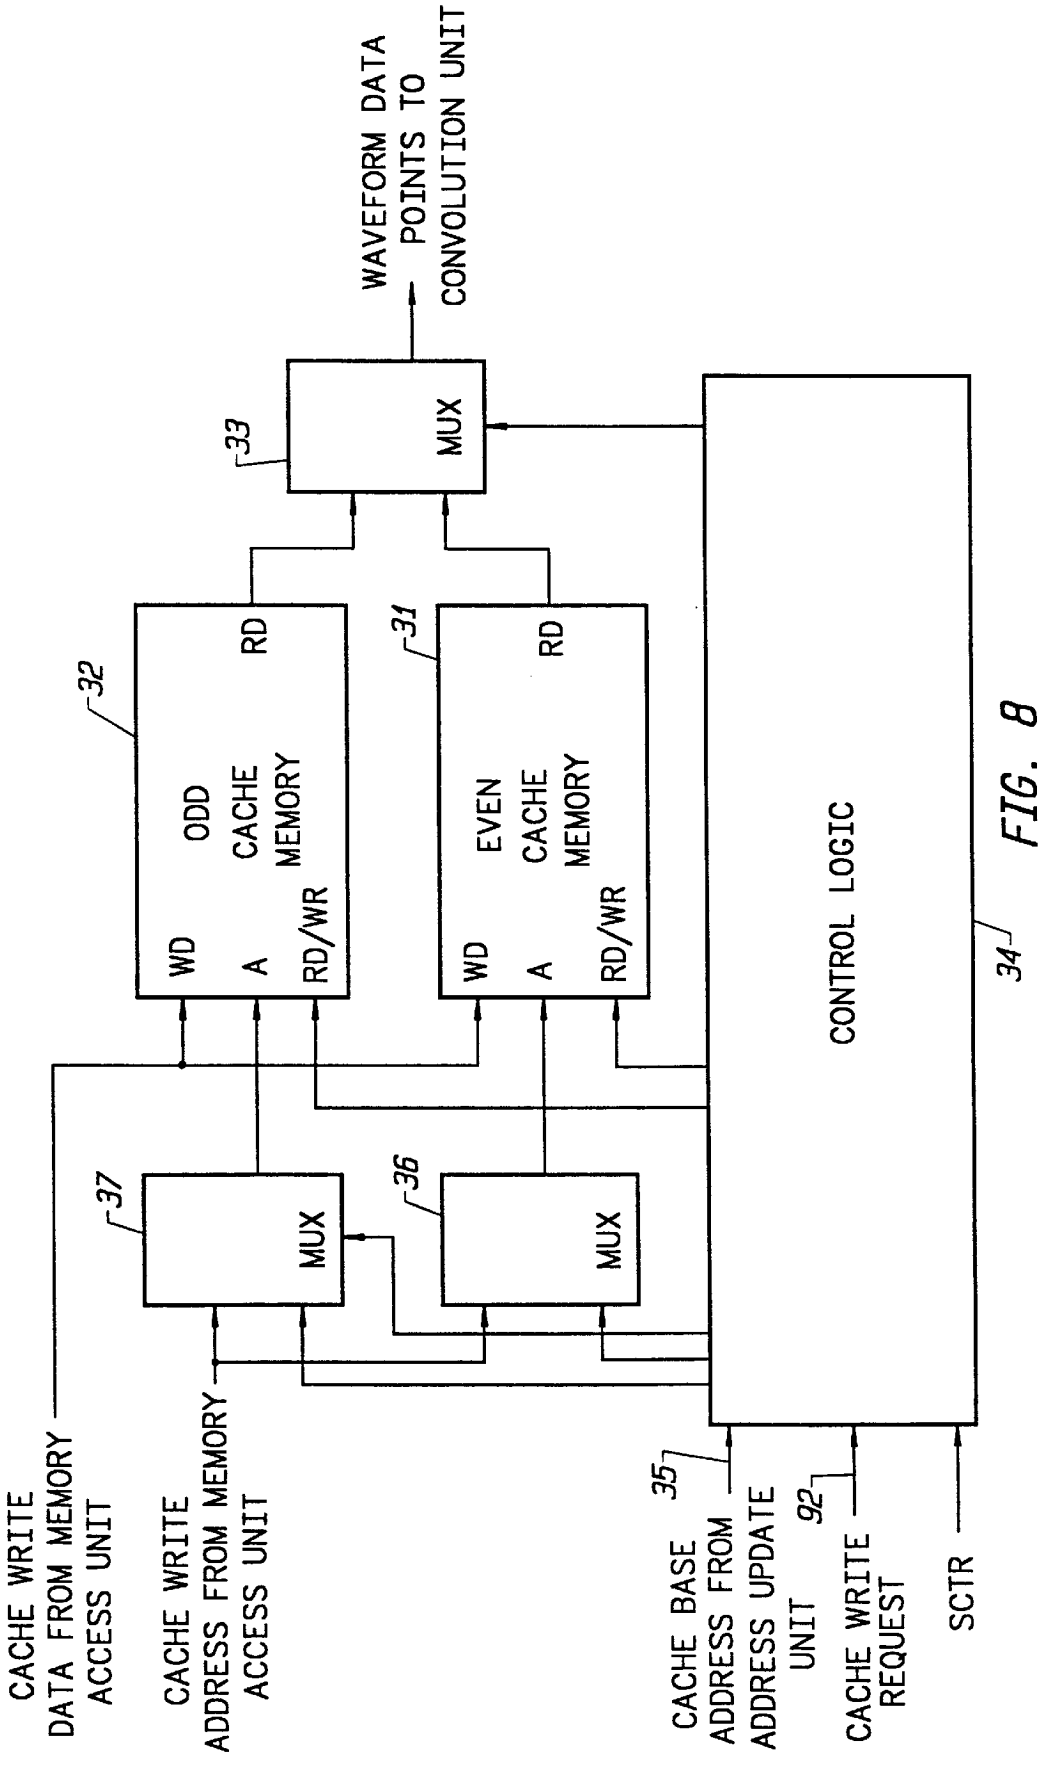
WAVEFORM DATA POINTS TO CONVOLUTION UNIT
33
MUX
32
ODD CACHE MEMORY
WD
A
RD/WR
RD
31
EVEN CACHE MEMORY
WD
A
RD/WR
RD
CONTROL LOGIC
34
37
MUX
36
MUX
CACHE WRITE DATA FROM MEMORY ACCESS UNIT
CACHE WRITE ADDRESS FROM MEMORY ACCESS UNIT
35
CACHE BASE ADDRESS FROM ADDRESS UPDATE UNIT
92
CACHE WRITE REQUEST
SCTR

FIG. 8

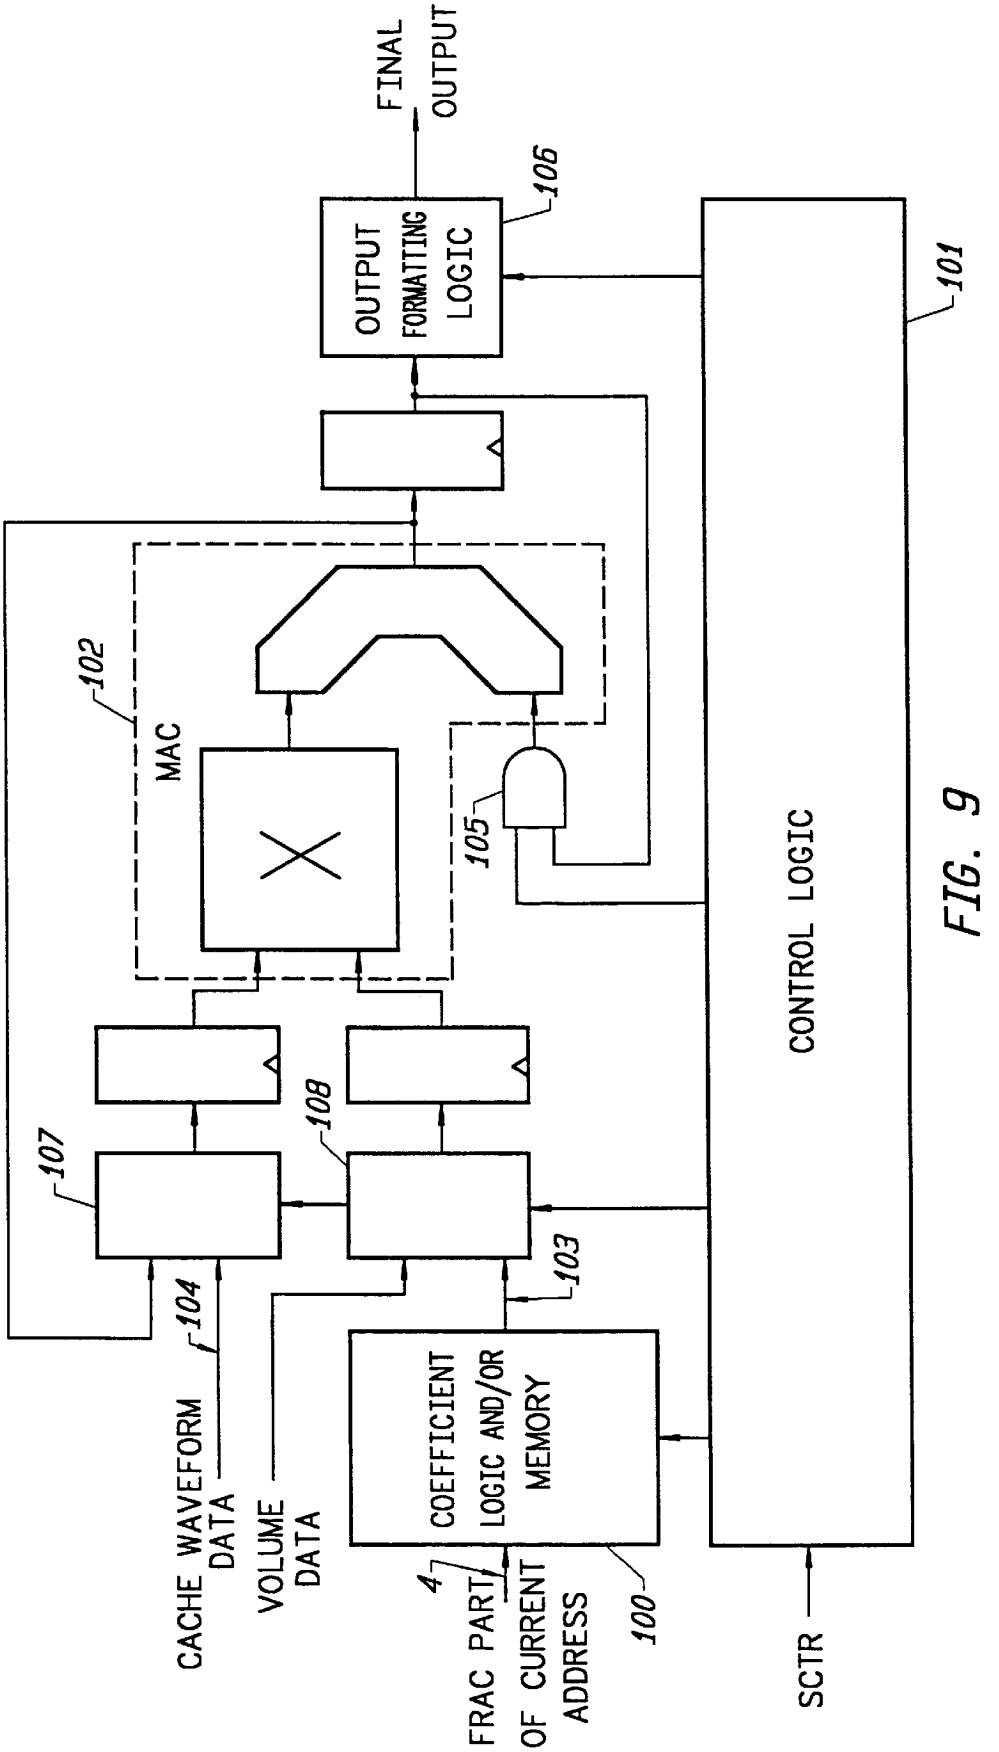
FINAL OUTPUT
OUTPUT FORMATTING LOGIC
106
101
102
MAC
105
CONTROL LOGIC
107
108
103
104
CACHE WAVEFORM DATA
VOLUME DATA
COEFFICIENT LOGIC AND/OR MEMORY
4
FRAC PART OF CURRENT ADDRESS
100
SCTR

FIG. 9

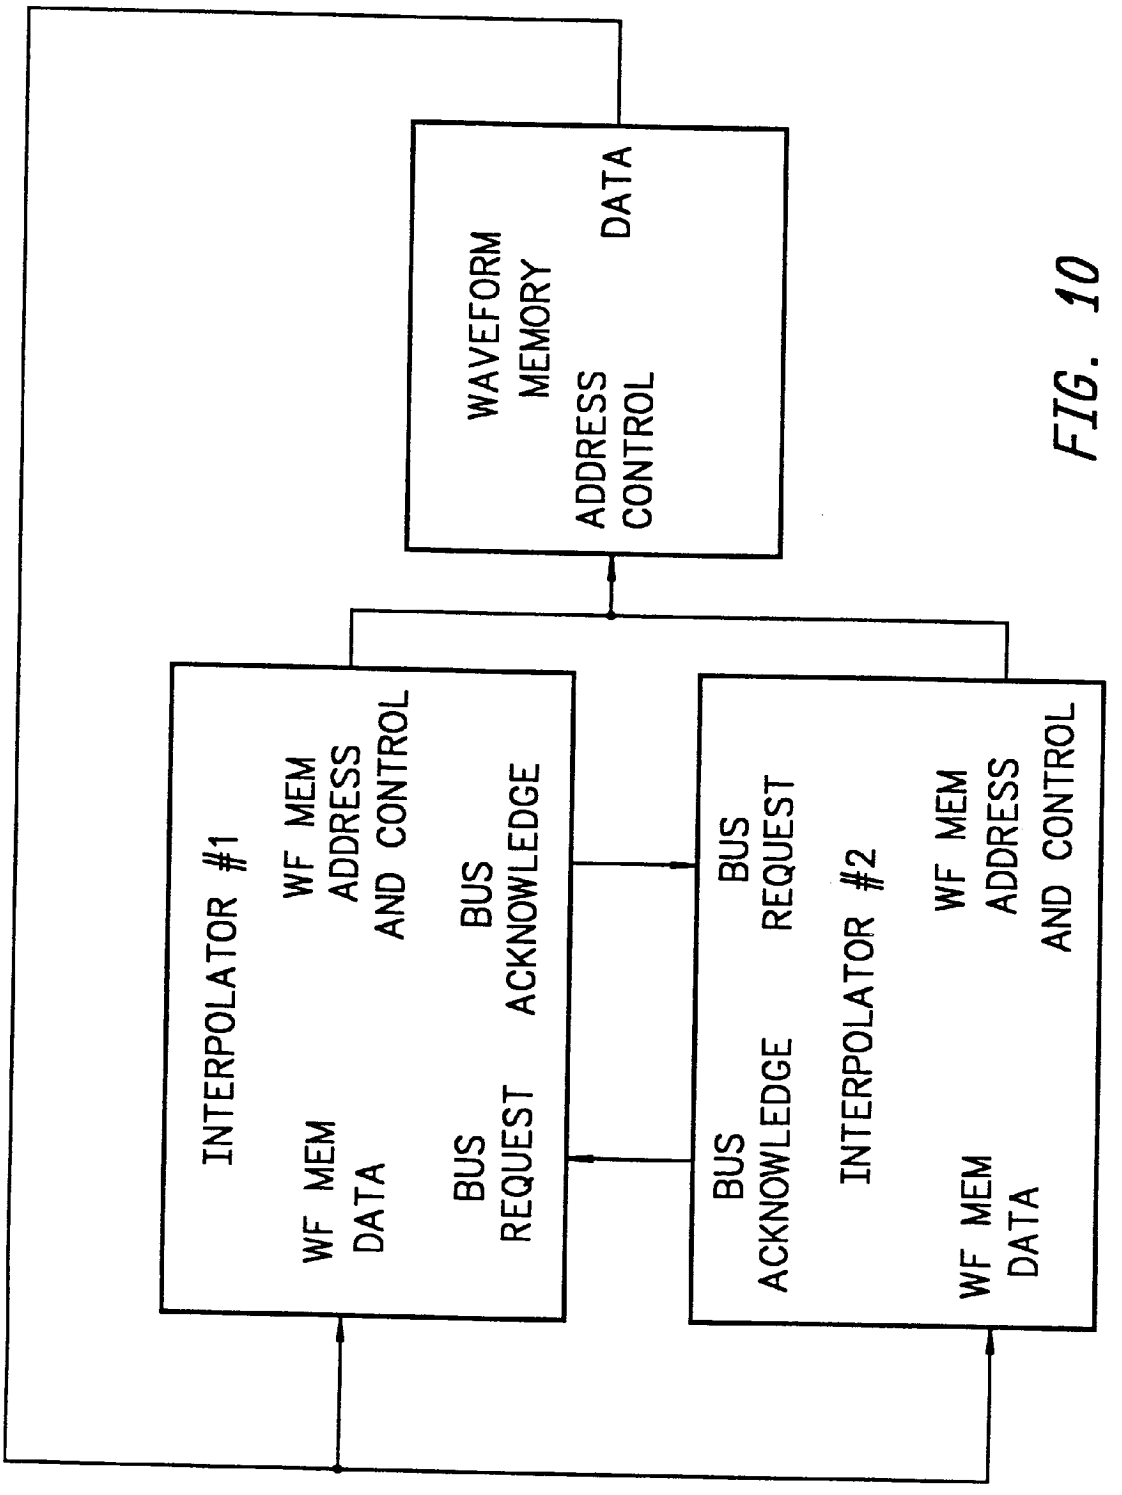
WAVEFORM MEMORY
ADDRESS
CONTROL
DATA
INTERPOLATOR #1
WF MEM ADDRESS AND CONTROL
BUS ACKNOWLEDGE
WF MEM DATA
BUS REQUEST
INTERPOLATOR #2
BUS REQUEST
WF MEM ADDRESS AND CONTROL
BUS ACKNOWLEDGE
WF MEM DATA

FIG. 10

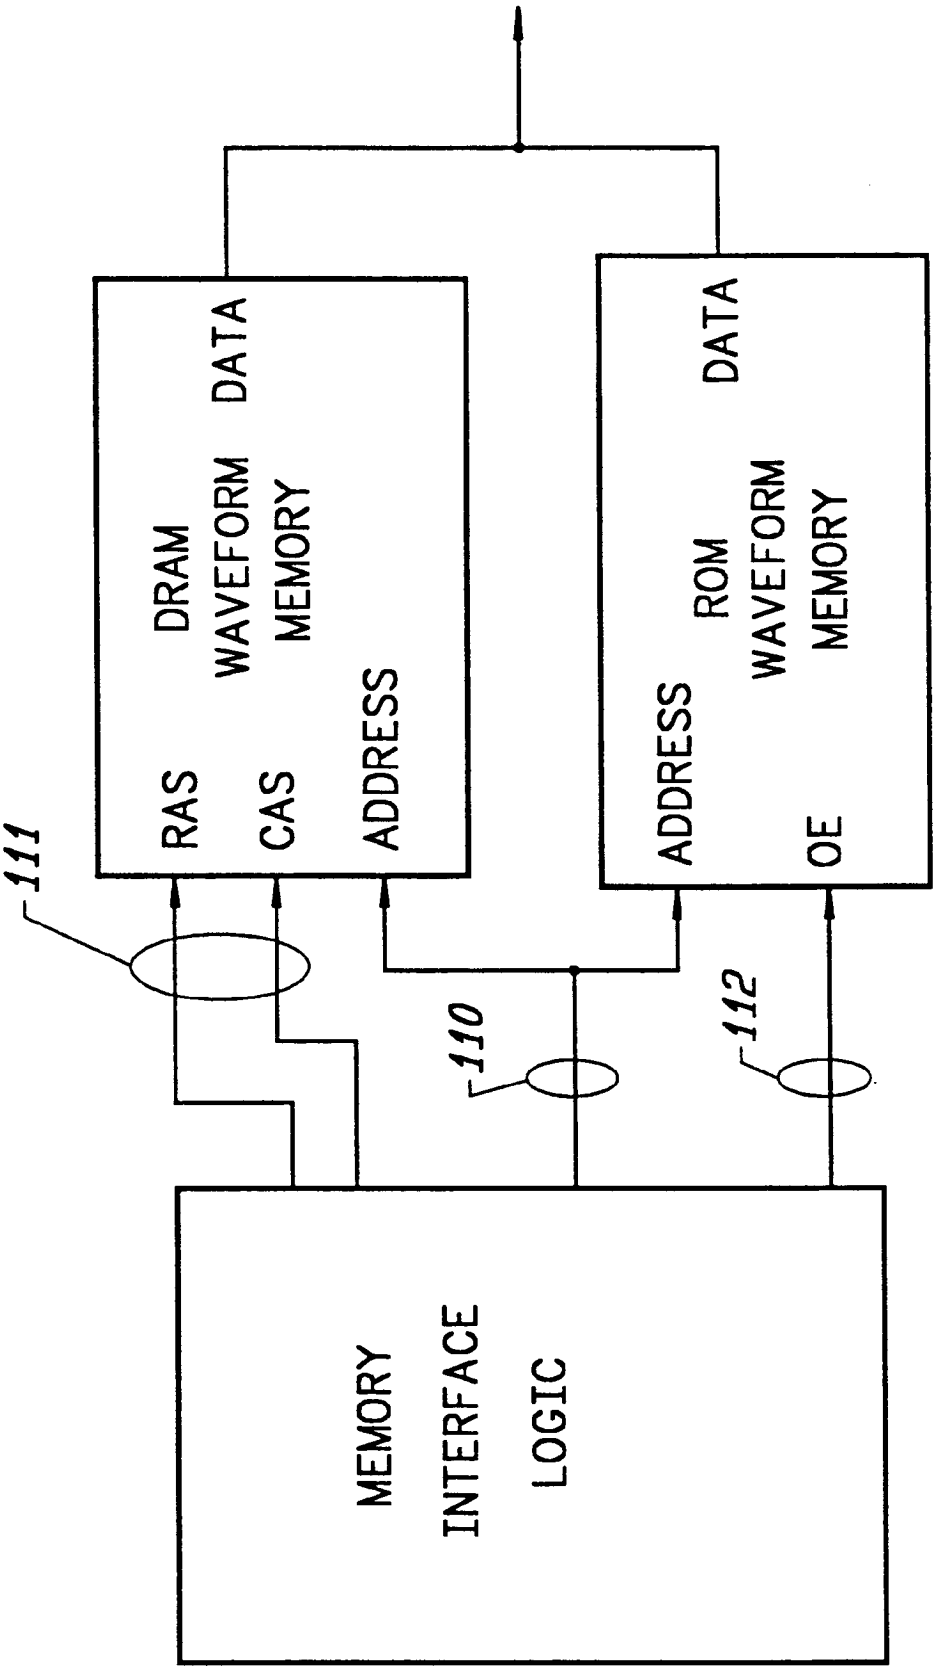
DRAM WAVEFORM MEMORY
RAS
CAS
ADDRESS
DATA
ROM WAVEFORM MEMORY
ADDRESS
OE
DATA
MEMORY INTERFACE LOGIC
111
110
112

FIG. 11

STATE COUNTER (8 PER CHANNEL)	0	1	2	3	4	5	6	7	0(NEXT CHANNEL)
CBA ADDRESS	NEW VALID CACHE BASE ADDRESS ----------							DON'T CARE	NEW VALID CBA
CACHE RAW READ OPERATION	NOP	READ POINT 0	READ POINT 1	READ POINT 2	READ POINT 3	READ POINT 4	READ POINT 5	READ POINT 6	NOP
CR READ OPERATION CBA EVEN	NOP	READ EVEN POINT	READ ODD POINT	READ EVEN POINT	READ ODD POINT	READ EVEN POINT	READ ODD POINT	READ EVEN POINT	NOP
CR READ OPERATION CBA ODD	NOP	READ ODD POINT	READ EVEN POINT	READ ODD POINT	READ EVEN POINT	READ ODD POINT	READ EVEN POINT	READ ODD POINT	NOP
EVEN CR OPERATION CBA EVEN	AVAIL FOR WR	READ POINT 0	AVAIL FOR WR	READ POINT 2	AVAIL FOR WR	READ POINT 4	AVAIL FOR WR	READ POINT 6	AVAIL FOR WR
ODD CR OPERATION CBA EVEN	AVAIL FOR WR	AVAIL FOR WR	READ POINT 1	AVAIL FOR WR	READ POINT 3	AVAIL FOR WR	READ POINT 5	AVAIL FOR WR	AVAIL FOR WR
EVEN CR OPERATION CBA ODD	AVAIL FOR WR	AVAIL FOR WR	READ POINT 1	AVAIL FOR WR	READ POINT 3	AVAIL FOR WR	READ POINT 5	AVAIL FOR WR	AVAIL FOR WR
ODD CR OPERATION CBA ODD	AVAIL FOR WR	READ POINT 0	AVAIL FOR WR	READ POINT 2	AVAIL FOR WR	READ POINT 4	AVAIL FOR WR	READ POINT 6	AVAIL FOR WR
MAU EVEN WRITE CBA EVEN	WRITE	WAIT	WRITE	WAIT	WRITE	WAIT	WRITE	WAIT	WRITE
MAU ODD WRITE CBA EVEN	WRITE	WRITE	WAIT	WRITE	WAIT	WRITE	WAIT	WRITE	WRITE
MAU EVEN WRITE CBA ODD	WRITE	WRITE	WAIT	WRITE	WAIT	WRITE	WAIT	WRITE	WRITE
MAU ODD WRITE CBA ODD	WRITE	WAIT	WRITE	WAIT	WRITE	WAIT	WRITE	WAIT	WRITE

FIG. 12

DIGITAL SAMPLING INSTRUMENT EMPLOYING CACHE MEMORY

CROSS-REFERENCE TO RELATED APPLICATIONS

The present application is a continuation of application Ser. No. 08/903,329, filed Jul. 29, 1997, U.S. Pat. No. 5,925,841, which is a division of application Ser. No. 08/636,827, filed Apr. 23, 1996, U.S. Pat. No. 5,698,803, which is a continuation of application Ser. No. 08/202,922, filed Feb. 28, 1994, abandoned, which is a division of application Ser. No. 07/882,178, filed May 11, 1992, U.S. Pat. No. 5,342,990, which is a continuation-in-part of application Ser. No. 07/462,392, filed Jan. 5, 1990, U.S. Pat. No. 5,111,727.

BACKGROUND OF THE INVENTION

The present invention may be efficiently implemented in a single VLSI integrated circuit of low cost. The present invention provides for a very high channel count, limited only be the speed and cost of the circuit and the average degree of upward pitch shifting required. Also, the present invention allows for use of multiple interpolator circuits to be used with a single waveform memory.

The present invention relates to electronic musical instruments, and more particularly to digital sampling instruments which create musical notes by reproducing recorded waveforms of musical instruments or sound effects, or mathematically calculated waveforms from a waveform memory at a variable playback rate. As has been previously disclosed in the above-identified cross-referenced patent application Ser. No. 07/462,392 filed Jan. 5, 1990 one technique for improving the performance of such instruments is the use of a cache memory and waveform interpolation. Such a technique increases the available channel count of the instrument by eliminating the waveform memory access time bottleneck which limits performance. While the basic use of cache memory has been previously described, there are several improvements beyond the preferred embodiment described in patent application Ser. No. 07/462,392 which are described herein.

SUMMARY OF THE INVENTION

It is an object of the present invention to provide improved N point interpolation variable pitch playback of musical sounds. The techniques so described can be efficiently implemented in the preferred embodiment by a single VLSI circuit of low cost. The preferred embodiment allows a very high channel count which will support many simultaneous musical notes. It also allows memory systems having variable access times, as well as the use of more than one interpolator circuit or chip with a single sound waveform memory.

Further described are techniques to optimize the performance of the system by decreasing the complexity of the VLSI circuit required to implement cache memory. Another technique described improves the cache system by increasing the degree to which upward pitch shifting on a minority of the channels can occur.

Other objects, features and advantages of the present invention will become apparent from the following detailed description when taken in conjunction with the accompanying drawings.

BRIEF DESCRIPTION OF THE DRAWINGS

The accompanying drawings, which are incorporated in and form a part of this specification, illustrate embodiments

of the invention and, together with the description, serve to explain, the principles of the invention:

FIG. **1** depicts a block diagram of an interpolator.

FIG. **2** depicts a block diagram of an interpolator with cache memory.

FIG. **3** depicts a block diagram of the present invention.

FIG. **4** depicts details of the address update unit of FIG. **3**.

FIG. **5** depicts details of the address register file of FIG. **3**.

FIG. **6** depicts details of the priority unit of FIG. **3**.

FIG. **7** depicts details of the memory access unit of FIG. **3**.

FIG. **8** depicts details of the cache memory unit of FIG. **3**.

FIG. **9** depicts details of the convolution unit of FIG. **3**.

FIG. **10** depicts a block diagram of a shared memory system.

FIG. **11** depicts a block diagram of the use of multiple waveform memory types.

FIG. **12** depicts a reservation table for cache RAM.

DETAILED DESCRIPTION OF THE PREFERRED EMBODIMENT

Reference will now be made in detail to the preferred embodiments of the invention, examples of which are illustrated in the accompanying drawings. While the invention will be described in conjunction with the preferred embodiments, it will be understood that they are not intended to limit the invention to those embodiments. On the contrary, the invention is intended to cover alternatives, modifications and equivalents, which may be included within the spirit and scope of the invention as defined by the appended claims.

The fundamental architecture of a sampling digital instrument which supports L channel pitch shifting by interpolation, as described in patent application Ser. No. 07/462,392 is shown in FIG. **1** and includes a waveform memory **1**, an address update unit **2** supplying an integer part **3** and a fractional part **4** of a current address for each of L channels, a memory address generator **5** producing waveform memory addresses **6** from the integer part **3** of the current address and a convolution unit **7** producing output samples depending on the fractional address **4** and the data **8** from waveform memory **1**. For each output sample and each of the L channels, the address update arithmetic circuit **10** provides a new current memory address consisting of an integer **3** and fractional **4** part which is created from the previous current address stored in register file **9** by adding a phase increment which is also stored in register file **9**, and the convolution unit's multiply-accumulator circuit **12** computes a sum of products of samples located in the waveform memory times coefficients (which depend on the fractional part of the memory address). This can be represented by the formulas:

$$A_n = A_{n-1} + P$$

$$A_n = A_{n-1} + P$$

-continued

$$Y_{i+f} = X_{i\frac{n-1}{2}} C_0(f) + X_{i\frac{n-3}{2}} C_1(f) + \ldots + X_i C_{\frac{n-1}{2}}(f) + \ldots + X_{i+\frac{n-1}{2}} C_n(f)$$

where An is the new current address, An-**1** is the previous current address, P is the phase increment, Yi+f is the output signal representing the current address with integer part i and fractional part f, Xm represents the waveform memory sample at address m, and Cn(f) represents the nth coefficient which is a function of f.

The coefficients Cn(f) are computed in the coefficient generator **11** either algorithmically from the fractional address f, by a table lookup based on f, or by a combination of these two approaches, such as linear interpolation of a limited size table. The output signals Y are digital output data **13** ready to be further processed by additional circuitry or by time domain multiplexing of some of the existing blocks.

The current memory address A for any given channel is continually being increased by the phase increment P at each output sample period. The magnitude of the phase increment determines the pitch shift from the original pitch, with an increment of unity being no shift, and increments smaller than unity shifting the pitch downward. As described in patent application Ser. No. 07/462,392, the phase increment is most commonly less than one. Consequently, the number of accesses of waveform memory can be reduced by the use of a cache memory. This will in turn increase the number of channels which can be supported by a single waveform memory.

A cache memory system is shown in FIG. **2**. In this system, not only does address update unit **2** produce an integer **3** and fractional **4** part of a new current address, it also produces a cache data required size (S) **14**, defined as:

$$S=Int(A_n)-Int(A_{n-1})$$

where Int(x) is the integer part of x.

The cache memory **15** contains M entries for each channel, and M must be an even power of two. A minimum of N entries is required for N point interpolation, that is M≧N. Thus, for example, linear interpolation (N=2) requires 2 or more entries per channel, and eight point interpolation M is greater than or equal to 8.

Cache memory **15** is written with S samples from waveform memory **1**, the last of which is located at the integer part **3** of the current memory address. The waveform memory addresses to accomplish this are generated by memory access unit **16**, which can operate independently and asynchronously from the address update and convolution units. Memory access unit **16** also generates the cache write address **17** at which the samples are stored.

Ultimately, the required samples **18** are read from the cache memory at sequential addresses, the last of which is located at the truncated least significant portion of the integer part **3** of the current memory address, and supplied to the convolution unit for multiplication and accumulator with coefficients derived from the fractional part **4** of the current address. This system operates according to the algorithm previously described in patent application Ser. No. 07/462,392.

The method described in patent application Ser. No. 07/462,392 has some limitations. Specifically, a dual port cache memory with separate write and read addresses is required, and the time required to complete all waveform memory accesses in the worst case is limited, thus placing an unnecessary limit on the upward pitch shifting capability of the circuit. The present invention eliminates both these limitations, and is shown in block diagram form in FIG. **3**.

As in all interpolation systems, an address update unit **2** produces a new current address for each output sample for each channel. However, in the present invention, the current address for each channel is stored in an external address register file **21**, along with the cache size and cache base address for each channel. Address register file **21** is only used by address update unit **2** for two cycles (one to fetch the old current address and another to store the new current address) for each channel, so it is available for other accesses which are controlled by memory access unit **23** on the remaining cycles. In the preferred embodiment, alternate cycles are available to the memory access unit, as determined by system state counter **26**.

Address update unit **2** produces a cache request signal **29** for a given channel if the integer part of the new current address differs from the integer part of the old current address. This request indicates that a waveform memory access is required. It also produces a cache size (S) and cache base address (CBA). The cache size is defined as above; the cache base address is a cumulating sum of the cache size, or:

$$CBA_n=(Int(A_n)-Int(A_{n-1})+CBA_{n-1}) \bmod M$$

where M is the cache memory size in entries per channel.

The need for a cache base address is twofold. First, the size of the cache is not an even power of two, the mechanism described previously in patent application Ser. No. 07/462,392 will not function because the mod operation in the equation above defining the CBA cannot be implicitly performed by truncation leaving only the least significant bits of the integer part of the current address. Thus a separate cache base address which can be modulo M must be maintained. However, even when M is an even power of two, there is reason to maintain a separate cache base address. A typical feature of digital sampling musical instruments is the ability to loop a sound in waveform memory. This is typically accomplished by an algorithm implemented in the address update unit in which the new current address is compared against a loop end address. If the end address has been exceeded, the size of the loop is subtracted to begin the loop of sound again near the start. This technique is well known to those skilled in the art. Because the size of the loop is not necessarily an integer multiple of the cache size, the cache address cannot be identical with the least significant bits of the current address if a loop is implemented. Instead, a separate cache base address must be maintained. When a loop occurs, the current address is updated according to a conventional looping algorithm, but the cache base address is simply increased modulo M by the phase increment as with an unlooped address update.

In FIG. **3**, priority unit **22** determines from the cache requests of all channels which channel's request should be honored first. It then supplies that priority channel's number **27** to memory access unit **23** and to address register file **21**. An idle signal **24** is also supplied to indicate that no requests are pending, and an accept signal **25** causes the priority unit to reset a channel's cache request as having been accepted for service.

Memory Access Unit **23** responds to a request for a given channel's service by asserting accept signal **25** after acquiring the channel number from priority unit **22**, and the current address, cache size, and cache base address for that channel from the address register file **21**. For each waveform memory location, beginning at the current address and

decrementing until S locations have been accessed, waveform memory **1** is accessed and the sample located at the indicated address is placed into cache memory **30**. The first fetched sample is placed at the cache base address for the channel, and subsequent samples, if any, are placed at successively decreasing locations module M.

The Convolution unit **7** behaves similarly to typical interpolation systems, responding to the fractional part **4** of the old current address to determine a set of N coefficients for N point interpolation. The old cache base address **28** is used to generate the starting location of the sample points in the cache memory, from which the sum of products of samples and coefficients is generated.

Master state counter **26** provides information used by all units as to the channel number under service and the processing state for that channel.

Now that the overall structure of the present invention has been explained, the details of each of the elements will be explained in detail.

The address update unit for the preferred embodiment shown in FIG. **4** is similar to that required for most interpolators, and should be familiar to those skilled in the art. Constants for control of the current address are stored in register file memory **51**, which is accessed according to address and read/write signals **60** and **61** supplied by control logic block **52**, which derives the register file access pattern from the master state counter. For cycles in which a register file access is not required, initial data can be written into the register file as requested by a controlling microprocessor. Constants accessed from register file **51** are loaded into register **53** for manipulation by the address update ALU **57**.

For example, the phase increment P is loaded into register **53** during a particular cycle. Simultaneously, data from the address register file is loaded into register **54**. Multiplexers **55** and **56** select appropriate fields of the data thus present in the registers to be manipulated by ALU **57** to produce the new current address and cache base address which are clocked into register **58**. Multiple cycles are typically required in order to provide for sound looping by comparison of the results of the incremented current address to a loop end address and conditional subtraction of a loop size. At some value of the state counter, the completely valid new current address and cache base address are available in register **58** and can be re-written into the address register file along with the cache size produced by cache request/size logic block **59**. This is repeated for each of the L channels to be processed according to the most significant portion of the state counter.

Other outputs from the address update unit are fractional part of the old current address and the old CBA. These are available from register **54** which has latched this data when accessed from the address register file prior to update. Another output is the cache request signal, which is derived from phase increment, old current address, old cache size, and any previously unhonored request from this channel, by cache request/size logic block **59**.

The cache size for the preferred embodiment is slightly more complex than the equation for S above due to two complicating factors. First, in the preferred embodiment, there are two levels of priority for cache requests, low and high, as explained below. It is possible for a low priority request to remain unhonored for an entire sample period. In this case, the unhonored request must continue to persist, and be added to the cache size as an additional cache entry to be fetched.

Furthermore, for reasons of channel startup, it is necessary that if the previous cache size was set to zero by the controlling processor, that any pending low priority request be ignored. Thus the logic for cache size follows the equation:

$$S=Int(A_n)-Int(A_{n-1})+(OldS!=0)*OldP$$

where OldS is the old cache size, and OldP is the state of the low priority bit for this channel from the previous sample period. Note that S must be saturated to never exceed M, the number of cache entries per channel.

A low priority cache request is generated if S is equal to one. A high priority cache request is generated if S is greater than one.

One skilled in the art will note that there is redundancy in the definitions of S and the low priority and high priority signals. In the particular case of M (the number of cache entries per channel) being an even power of two, and N (the number of interpolation points) being M−1, the required values of S are zero to an even power of two, which must be very inefficiently encoded in log2M+1 bits. Encoding S=0 indicating zero cache entries required, S=1 to indicate either 1 or 2 cache entries required depending on the priority (low=1, high=2) and S>1 indicates S+1 entries required reduces the bit requirement to an efficient log2M bits.

The address register file shown in FIG. **5** is once again a block that should be understood by those skilled in the art. Register file memory **61** contains a location for each channel, in which is stored the current address, cache size, and cache base address on alternate cycles as based on the master state counter, access to the memory is given to the memory access unit by reading the address indicated by the channel number provided by the priority logic. The remaining cycles can be allocated to reading and subsequently writing the data for the channel under service by the address update unit, whose channel number is determined from the most significant part of the state counter, or to writing initialization data into the register for memory from a controlling processor. The selection of the write data from controlling processor or address update unit is performed by multiplexer **63**, which along with register file address and read/write is controlled by control logic block **62** which derives its sequence from the master state counter.

The priority unit is shown in detail in FIG. **6**. Before the function of the priority unit can be explained, some further clarification of the function of this circuit is necessary.

While the minimum size for a cache memory is M=N entries per channel for an N point interpolator, if the cache memory is made larger, for example N+1 locations per channel, a significant benefit can be achieved.

Consider the typical operation of the cache memory and address update unit. The address update unit will determine the current address (both integer and fractional part) and initiate the convolution cycle for the corresponding output point. Note that this is why the old fractional part and CBA are used by the convolution unit and cache memory control circuit. The address update unit will then add the phase increment to the current address to produce a new address. If the new address has changed in integer part from the old address (S!=0), the memory access unit is requested to fetch the data from the new address and any intervening addresses into the cache. This must be done before the convolution cycle for the next output point, if only locations are present in the cache, because all N must be valid for an N point interpolator.

Since each channel is operating with a phase increment of arbitrary fractional part, the number of waveform memory cycles will jitter between two values for each channel. For example, if the phase increment were exactly 0.5, alternate

cycles would require zero or one waveform memory access. The typical case would involve a much more complex pattern of accesses. However, it is clear that there is a "worst case" instance in which every channel would have its larger number of waveform memory accesses required.

If the memory access unit were not capable of servicing all of these accesses, the cache would not be properly filled for some channel and an incorrect convolution would result. Thus the maximum upward pitch shift capability for such a unit with N locations per channel for an N point interpolator, L channels, and M waveform memory accesses per sample period, will be such that the sum of the integer parts of the phase increments for all the channels is less than M-L. Thus, for example, if M equals 64 and L equals 64, then no pitch shift can exceed unity.

The above situation can be alleviated by the use of an oversized cache memory and a priority circuit. If the cache memory has one more location than the number of points of interpolation, then one cache memory location can remain invalid for more than one cycle. Depending on the phase increment for that particular channel, the waveform memory access for that channel can be postponed for one or more cycles while accesses which are required more quickly are performed. It can be shown that in this case, the sum of the phase increments themselves cannot exceed M. This is a considerably better situation than that above.

The priority circuit must thus distinguish between two types of memory access unit requests, those which can be serviced later, and those which must be serviced before the next cycle. Within each category, a priority is established in which as a channel gets closer to its next service by the address update unit, it gets higher priority. Such a circuit is shown in FIG. **6**.

There are two cache service request lines, one of low and one of high priority. The generation of these by the address logic has been explained above. A shift register **70** is enabled once for each channel, and thus maintains a bit for each channel for each of these two priority levels. The 2N outputs of the shift registers are routed into a standard priority encoder **71**, whose output **72** is the number of the highest active input of high priority, if any, or if none, then the number of the highest active input of low priority. If no inputs are active, an idle flag **73** is asserted. The priority channel and idle flag are latched in register **74**. Because the requests are being shifted within the shift register for each new channel processed, the number in register **74** must be added to the current active channel number **75** as provided by the master state counter to give the current priority channel. This sum at the output of adder **76** is stored in enabled register **77**, and supplied to the address file as the memory access unit channel.

When enabled register **77** becomes enabled with non-idle data, a new priority channel has been accepted for service, and that channel's priority can thus be reset. This is determined by gate **78** whose output becomes active only if the output of register **74** does not correspond to an idle priority, and either the output of enable register **77** is idle, or the memory access unit is acknowledging acceptance of the channel indicated by enabled register **77**. In this case, decoder **79** is enabled and resets both shift register bits corresponding to the channel which is indicated by register **74**. Enabled register **77** is disabled after it has accepted the new channel because the memory access unit will negate its accept line and is re-enabled when the memory access unit has completed its processing of the previous channel and hence asserts the accept signal again.

FIG. **7** shows the details of the memory access unit. When a new priority channel **80** becomes valid from the priority unit as indicated by idle signal **81** becoming low, these signals along with the current address integer part, cache size, and cache base address for the channel as supplied by the address register file are each latched into enabled register **82** which is enable during cycles when the memory access unit has had an access to the address register file as determined by control logic **83** from the master state counter. AND gate **84** then accepts the data from enabled register **82**, and initiates a set of memory access cycles if the data in register **82** is valid (not idle) and the memory access unit is not already busy. Control logic **83** begins a cycle by presetting down counters **85**, **86**, and **87** with the current address integer part, cache base address, and cache size respectively, and indicates the memory access unit is busy by asserting line **88**. If the cache size is zero, the cycle is aborted. This is necessary for the startup condition on a channel so that the controlling processor, by setting cache size to zero, can halt any activity which will fill the cache.

If the cycle is not aborted, the memory interface logic **89** determines from the current address what type of memory is being used in this portion of the address space, and initiates an appropriate memory cycle by providing properly formatted address and control signals on waveform memory address and control lines **90**.

FIG. **11** shows a typical connection for multiple memory types. Note that some address and control lines can be shared by both types of memory (in this case, DRAM and ROM memory have been chosen) as illustrated by signals of group **110**, while others are dedicated to either one or the other memory, such as group **111** for the DRAM, and **112** for the ROM. When the waveform memory data is present on the memory data bus **8** of FIG. **7**, it is latched along with the channel number and the CBA counter value in enabled register **91**, and memory interface logic **89** indicates to control logic **83** that the current memory cycle is complete. Control logic **83** then asserts cache write request line **92**, indicating that the cache write data and cache address present in enabled register **91** have become valid, and also causes all three counters **85**, **86**, and **87** to be decremented. If the cache size in counter **87** has become zero, the request is complete and the busy line is brought inactive, enabling another request. If the cache size is non-zero, another memory cycle is initiated with the decremented current address and CBA.

Note that dynamic memory can use page mode accesses for multiple memory cycles to adjacent memory locations, and that the logic to accomplish this is straightforwardly included into the memory interface logic. Also note that the asynchronous nature of the memory interface logic allows for optionally including logic for bus arbitration signals **93** and **94** (which would typically be a bus acknowledge and bus request respectively) which would allow sharing of a single waveform memory among multiple interpolator systems or chips. Naturally in such a case, the waveform memory address and control signals **90** would have to be capable of being output disabled by memory interface logic **89** when it was not the bus master. Such a multiple interpolator system is shown in FIG. **10**.

Details of the cache memory unit are shown in FIG. **8**. Note that the silicon area required to implement dual port memory is nearly twice that require for single port memory. Hence it is cost effective to utilize single port memory whenever possible. Avoiding a dual port cache memory is accomplished by noting that the convolution algorithm requires the sum of products, each of which comes from successive locations in the cache memory. In other words, for an N point interpolator, the output will be the sum of a first coefficient times a first point in the cache memory, plus a second coefficient times the next point in the cache memory, et cetera, up to the Nth coefficient times the final sequential cache memory datum. Consequently, the cache memory can be divided into two separate single port memories as shown in FIG. **8**. Even cache memory **31** contains the even locations and odd cache memory **32** contains the odd

locations. Multiplexer **33** routes the output data from the cache memory active for reading data to the convolution unit. The convolution unit will thus access alternate memories, leaving the unaccessed memory free for accepting data from the waveform memory as directed by the memory access unit.

Control logic **34** determines from the read cache base address **35** supplied by the address update unit and the master state counter which memory will be active for reading during a given cycle.

This determines the state of multiplexer **33** as well as that of address multiplexers **36** and **37**, which determine if memory **31** or **32** respectively is getting its address from the memory access unit or from the cache base address and channel number under service by the convolution unit.

By allocating at least one cycle of the convolution circuit to a function other than the convolution itself, either an idle cycle or one for scaling the volume of the resulting sum of products, one cache memory cycle per output point can require no read cycle and thus have both cache memories available for writing memory access unit data. As shown in the reservation table in table 1 (see FIG. **12**) this then guarantees that the proper cache memory is available for a write access from the memory access unit within no more than two cycles. At typical clock rates, this occurs faster than the access times of memories suitable for waveform storage. Control logic **34** thus responds to a cache write request on line **92** by writing the waveform memory data at the required cache memory address within two cycles.

The details of the convolution unit are shown in FIG. **9**. Coefficient logic **100** derives the N coefficients as a function of current address fractional part **4** according to the master state counter as interpreted by control logic **101**. This can be accomplished according to various algorithms, some of which are described in patent application Ser. No. 07/462,392. Multiply/accumulator **102** forms the sum of products required by the interpolation algorithm from the coefficients **103** times the cache memory data **104**. Control logic **101** can reset the cumulating sum in the accumulator using AND gates **105** at the beginning of each new channel. The ultimate sum of products is then formatted, for example, into a serial data stream for further processing, by output formatting logic **106**.

Provision is also made, by use of multiplexers **107** and **108**, to use any idle cycles of the multiply/accumulator (one of which is necessary if single port cache memory has been used as described above) for alternate functions. For example, the output of the multiply/accumulator's final sum can be routed into one input of the multiplier by multiplexer **107**, while a volume scaling number is routed into the other input by multiplexer **108**, allowing the idle cycle to provide an individual volume control for each channel before the final sum of products is routed to output formatting logic **106**.

The foregoing descriptions of specific embodiments of the present invention have been presented for purposes of illustration and description. They are not intended to be exhaustive or to limit the invention to the precise forms disclosed, and it should be understood that many modifications and variations are possible in light of the above teaching. The embodiments were chosen and described in order to best explain the principles of the invention and its practical application, to thereby enable others skilled in the art to best utilize the invention and various embodiments with various modifications as are suited to the particular use contemplated. It is intended that the scope of the invention be defined by the claims appended hereto and their equivalents.

What is claimed is:

1. A digital sampling instrument for the multichannel interpolative playback of digital audio data samples stored in a waveform memory, comprising:

a cache memory storing two or more waveform memory samples for each channel;

control logic to access two adjacent ones of said waveform memory samples from said cache memory; and

an interpolator configured to linearly interpolate between said two adjacent waveform memory samples to form a linear interpolation result.

2. An instrument as in claim **1** wherein said linear interpolation is mathematically equivalent of the formula: $Y_{i+f}=X_{i-1}*f+X_{i-2}*(1-f)$ wherein X_{i-1} and X_{i-2} are said adjacent waveform memory samples, f is a coefficient, and Y_{i+f} is said result of said linear interpolation for each channel.

3. An instrument as in claim **2** wherein said coefficient f is the fractional part of a memory address.

4. An instrument as in claim **3** wherein said memory address is computed by repeated addition of an increment with a fractional part to a base address.

5. An instrument as in claim **4** wherein said waveform memory is implemented in dynamic RAM.

6. An instrument as in claim **1** wherein the linear interpolation result is at a certain sample rate and wherein one or more of said digital audio data samples are at a sample rate below said certain sample rate.

7. A method for the multichannel interpolative playback of digital audio data samples stored in a waveform memory, comprising:

storing two or more waveform memory samples for each channel in a cache memory;

accessing two adjacent ones of said waveform memory samples from said cache memory; and

linearly interpolating between said two adjacent waveform memory samples to form a linear interpolation result.

8. A method as in claim **7** wherein said linear interpolation is mathematically equivalent of the formula: $Y_{i+f}=X_{i-1}*f+X_{i-2}*(1-f)$ wherein X_{i-1} and X_{i-2} are said adjacent waveform memory samples, f is a coefficient, and Y_{1+f} is said result of said linear interpolation for each channel.

9. A method as in claim **8** wherein said coefficient f is the fractional part of a memory address.

10. A system for the multichannel interpolative playback of digital audio data samples, comprising:

a waveform memory storing said digital audio data samples;

a bus coupled to said waveform memory;

a digital sampling instrument coupled to said bus, said instrument including

a cache memory storing two or more waveform memory samples for each channel;

control logic to access two adjacent ones of said waveform memory samples from said cache memory; and

an interpolator configured to linearly interpolate between said two adjacent waveform memory samples to form a linear interpolation result.

11. A system as in claim **10** wherein said linear interpolation is mathematically equivalent of the formula: $Y_{i+f}=X_{i-2}*X_{i-2}*(1-f)$ wherein X_{i-1} and X_{i-2} are said adjacent waveform memory samples, f is a coefficient, and Y_{i+f} is said result of said linear interpolation for each channel.

12. A system as in claim **11** wherein said coefficient f is the fractional part of a memory address.

* * * * *